

# DSO 545: Statistical Computing and Data Visualization

*Abbass Al Sharif*

*Fall 2019*

## Lab 7: Data Wrangling

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

### 1. Data Science Tools: map(), apply(), and lambda functions

1. Load the `titanic.csv` dataset into a pandas DataFrame.

```
data = pd.read_csv("titanic.csv")
```

2. Use the `map()` function to create a new variable call “Sex\_Numeric” (0 for males, and 1 for females).

```
data["Sex_Numeric"] = data['Sex'].map({"female":1, "male":0})
```

3. What is the percentage of females in the Titanic dataset?

```
data["Sex_Numeric"].mean()
```

### OR

```
## 0.35241301907968575
```

```
data.Sex.value_counts()/data.shape[0]
```

```
## male      0.647587
```

```
## female    0.352413
```

```
## Name: Sex, dtype: float64
```

4. Create a new variable called “Fare\_ceil” to show the fare value of the trip rounded up.

```
data['Fare_ceil'] = data['Fare'].apply(np.ceil)
```

5. Load the `drinks.csv` dataset into a dataset called `drinks`.

```
drinks = pd.read_csv("drinks.csv")
```

6. Find the max beer, spirit, and wine servings among all countries.

```
drinks.loc[:, 'beer_servings': 'wine_servings'].apply(max, axis = 0)
```

```
## beer_servings      376
```

```
## spirit_servings     438
```

```
## wine_servings      370
```

```
## dtype: int64
```

7. Find the max serving among beer, spirit, and wine for each country.

```
drinks.loc[:, 'beer_servings': 'wine_servings'].apply(max, axis = 1).head()
```

```
## 0      0
```

```
## 1     132
```

```
## 2      25
## 3     312
## 4     217
## dtype: int64
```

8. Find the max category of serving among beer, spirit, and wine for each country.

```
drinks.loc[:, 'beer_servings': 'wine_servings'].apply(np.argmax, axis = 1).head()
```

```
## 0      beer_servings
## 1    spirit_servings
## 2      beer_servings
## 3      wine_servings
## 4      beer_servings
## dtype: object
```

```
##
```

```
## /anaconda3/envs/r-reticulate/lib/python3.7/site-packages/numpy/core/fromnumeric.py:61: FutureWarning
## The current behaviour of 'Series.argmax' is deprecated, use 'idxmax'
## instead.
```

```
## The behavior of 'argmax' will be corrected to return the positional
## maximum in the future. For now, use 'series.values.argmax' or
## 'np.argmax(np.array(values))' to get the position of the maximum
## row.
```

```
## return bound(*args, **kwds)
```

9. Create a function that takes an input value x and returns it four-fold.

```
def four_fold(x):
    return 4*x
```

```
four_fold(3)
```

```
## 12
```

10. Create a new column in the drinks dataset `beer_4fold` which multiplies the value of the beer servings in each country by 4.

```
drinks['beer_4fold'] = drinks['beer_servings'].apply(four_fold).head()
```

11. Use a lambda function to answer the previous question, i.e. don't use the function `four_fold()`.

```
drinks['beer_4fold'] = drinks['beer_servings'].apply(lambda x:4*x).head()
```

```
## OR
```

```
drinks['beer_4fold'] = list(map(lambda x:4*x, drinks['beer_servings']))
```

## 2. Data Wrangling

### Filtering Data

12. Find all passengers who are above 30 years old.

```
s1 = data[data.Age > 30]
```

13. Find all female passengers who are above 30 years old.

```
s2 = data[(data.Age > 30) & (data.Sex == "female")]
```

## Selecting Variables

14. Create a dataframe which has only two columns: PassengerId, Survived, and Cabin.

```
s3 = data[['PassengerId', 'Survived', 'Cabin']]
```

15. Create a dataframe that has all variables in the dataset except the “Cabin” variable.

```
s4 = s3.drop('Cabin', axis = 1)
```

## Arranging Data

16. Find all female passangers who are above 30 years old. The resulting dataframe should have three columns: PassengerId, Survived, Age. The dataframe should be arranged by age in descending order.

```
data.loc[(data.Age > 30) & (data.Sex == "female"),  
         ["PassengerId", "Survived", "Age"]].sort_values('Age', ascending = False)
```

```
##      PassengerId  Survived  Age  
## 275           276         1  63.0  
## 483           484         1  63.0  
## 829           830         1  62.0  
## 366           367         1  60.0  
## 195           196         1  58.0  
## ..           ...         ...   ...  
## 215           216         1  31.0  
## 328           329         1  31.0  
## 18             19         0  31.0  
## 318           319         1  31.0  
## 767           768         0  30.5  
##  
## [103 rows x 3 columns]
```

## Grouping and Summarizing Data

17. Find the average age of both male and female passangers.

```
data.groupby('Sex')['Age'].mean()
```

```
## Sex  
## female    27.915709  
## male      30.726645  
## Name: Age, dtype: float64
```

18. Find the median fare for passengers according to their class.

```
data.groupby('Pclass').Fare.median()
```

```
## Pclass  
## 1      60.2875  
## 2      14.2500  
## 3       8.0500
```

```
## Name: Fare, dtype: float64
```

19. Find the average and median, and the difference between mean and median age of both male and female passengers.

```
data.groupby('Sex')['Age'].agg({"Average": 'mean',  
                               "Median": 'median',  
                               "Diff": lambda x: x.mean() - x.median()})
```

```
##           Average  Median    Diff  
## Sex  
## female  27.915709    27.0  0.915709  
## male    30.726645    29.0  1.726645  
##
```

```
## /anaconda3/envs/r-reticulate/bin/python:3: FutureWarning: using a dict on a Series for aggregation  
## is deprecated and will be removed in a future version. Use                named aggregation instead
```

```
## >>> grouper.agg(name_1=func_1, name_2=func_2)
```

20. Find the average age for males and females who survived the Titanic disaster.

```
data.groupby(['Sex', 'Survived'])['Age'].mean()
```

```
## Sex    Survived  
## female 0         25.046875  
##        1         28.847716  
## male   0         31.618056  
##        1         27.276022  
## Name: Age, dtype: float64
```

```
# we can reset the index to make the data a dataframe
```

```
data.groupby(['Sex', 'Survived'])['Age'].mean().reset_index()
```

```
##      Sex  Survived    Age  
## 0  female         0  25.046875  
## 1  female         1  28.847716  
## 2   male         0  31.618056  
## 3   male         1  27.276022
```

21. Find the median fare for passengers embarked from different ports and among different classes.

```
data.groupby(['Pclass', 'Embarked']).Fare.median()
```

```
## Pclass  Embarked  
## 1       C         78.2667  
##        Q         90.0000  
##        S         52.0000  
## 2       C         24.0000  
##        Q         12.3500  
##        S         13.5000  
## 3       C          7.8958  
##        Q          7.7500  
##        S          8.0500  
## Name: Fare, dtype: float64
```

## Mutating Data

22. Create a new column (age\_cat) in the dataset to based on the age variable (“young” if age <=50 otherwise “older”).

```
data['age_cat'] = data['Age'].apply(lambda x: "Young" if x<=50 else "Older")
data.age_cat.unique()
```

```
## array(['Young', 'Older'], dtype=object)
```

23. Create a new column (age\_cat1) in the dataset to based on the age variable (“young” if age <=20, “mature” if 20<age<=50 otherwise “older”).

```
data['age_cat1'] = data['Age'].apply(lambda x: "Young" if x <= 20 else ("Mature" if x <= 50 else "Older")
data.age_cat1.unique()
```

```
## array(['Mature', 'Older', 'Young'], dtype=object)
```

24. Create the following Tree map using the mpg.csv dataset and squarify Python package.

